

Runbook — Confluent Platform + Flink on minikube

End-to-end operational guide for running `confluent-kafka-isotope` locally on **minikube**: cluster → Kafka/Confluent Platform → Flink → SQL reports → traffic → observe → teardown. Every step maps to a target in the [Makefile](#), which is the source of truth.

This is the **Confluent Platform + Flink (self-managed) on minikube** path. The Confluent Cloud for Apache Flink (CCAF) path is entirely different — Terraform-driven via `make cc-flink-reports-up`, no minikube — see [root README §3.3](#).

Table of Contents

- [1.0 Prerequisites \(once per machine\)](#)
- [2.0 Cluster + Confluent Platform](#)
- [3.0 Flink](#)
- [4.0 Port-forward Kafka](#)
- [5.0 Deploy the 7 Flink reports \(as a CMF Application\)](#)
- [6.0 Drive traffic \(required to see report rows\)](#)
- [7.0 Observe](#)
 - [7.1 \[OPTIONAL\] Showcase the two provenance collectors](#)
 - [7.1.1 Fan-in \(Merge\) Provenance — Three things to point at](#)
 - [7.1.1.1 The merged record carries a *fresh* trace, not a stolen one](#)
 - [7.1.1.2 The merge trace ID is derived, not minted — the window is readable inside it](#)
 - [7.1.1.3 Both reconciliations hold exactly](#)
 - [7.1.2 State Provenance — Three things to point at](#)
 - [7.1.2.1 The version chain, inline](#)
 - [7.1.2.2 The whole topic is internally consistent](#)
 - [7.1.2.3 No hops anywhere](#)
 - [7.1.3 The punchline — one order, three registers](#)
- [8.0 Teardown](#)
- [9.0 Minimal path \(no Flink\)](#)
- [10.0 Troubleshooting](#)

1.0 Prerequisites (once per machine)

```
make install-prereqs    # docker, kubectl, minikube, helm, gettext,
gradle, openjdk17
make check-prereqs      # verify they're on PATH
```

Default minikube sizing (override via env): `MINIKUBE_CPUS=6`, `MINIKUBE_MEM=20480`, `MINIKUBE_DISK=50g`. The node architecture is auto-detected so the right `cp-flink` image (amd64/arm64) is selected.

2.0 Cluster + Confluent Platform

```
make cp-up          # = check-prereqs → minikube-start → operator-
install → cp-deploy
make cp-watch       # watch pods come up (Ctrl+C to exit); or: make
cp-status
```

cp-up boots minikube, installs the CFK (Confluent for Kubernetes) operator, and deploys Kafka (KRaft) + Schema Registry + Connect + ksqlDB + REST Proxy + Control Center into the **confluent** namespace.

cp-up deliberately does **not** bring up Flink — run §3.0 Flink separately.

3.0 Flink

```
make cp-flink-up      # cert-manager → operator → MinIO → CMF 2.4 →
env → RBAC → app image
make cmf-status       # (~5 min the first time)
```

This installs cert-manager, the Confluent Flink Kubernetes Operator, **MinIO** (S3-compatible store for CMF artifacts), **CMF 2.4** (with artifact management + the writable environment catalog enabled), creates the **dev-local** Flink environment, applies the Flink RBAC, and builds the custom **cp-flink** 2.1 image (**isotope-cp-flink-sql:local**) that bakes in the Kafka + Avro SQL connectors and the S3 filesystem plugin.

The reports run as a CMF *Application*, not SQL statements. CMF's SQL-statement runtime (**io.confluent.flink.FlinkCompiledPlanExecutor**) ships only in the **cp-flink-sql** image, which exists only at **Flink 1.19** — and two reports are **ProcessTableFunctions**, a **Flink 2.x** feature. So all seven reports deploy as a single Flink 2.1 CMF **Application** (entry point **IsotopeReportsJob**), which runs the same **.fql** and surfaces in CMF / Control Center as a managed application. No raw session cluster is created — if you want one for ad-hoc SQL, deploy it separately with **make flink-deploy** (see §7.0 Observe).

CMF license. The **cp-cmf** image ships with a date-locked trial license. Once it expires, the CMF pod **CrashLoopBackOffs** on startup (**LicenseInitiator** throws) and **make cp-flink-up** fails at the CMF readiness wait. To run past expiry, supply a Confluent license via a secret and point **CMF_LICENSE_SECRET** at it:

```
kubectl create secret generic confluent-license-for-cmf -n confluent \
  --from-file=license.txt=/path/to/license.txt # key MUST be
license.txt
make cp-flink-up CMF_LICENSE_SECRET=confluent-license-for-cmf
```

Export **CMF_LICENSE_SECRET** in your shell to make plain **make cp-flink-up** pick it up.

4.0 Port-forward Kafka

```
make kafka-pf-up          # localhost:30092 → Kafka, localhost:8081 → SR
```

Prereq for everything the host-run gradle app does — `App.java`'s defaults already point at `localhost:30092` / `localhost:8081`. Stop with `make kafka-pf-down`.

5.0 Deploy the 7 Flink reports (as a CMF Application)

```
make cp-flink-reports-up  # builds the app shadow JAR (:ptf:shadowJar) →
pre-creates                # 4 source + 7 sink topics → uploads the JAR as
a cmf://                    # artifact (MinIO) → deploys the
FlinkApplication → waits RUNNING
```

All seven reports run in **one** Flink 2.1 CMF Application (`isotope-reports`): five pure Flink SQL plus two JAR-backed `ProcessTableFunctions` (`LatencyPercentilesPTF`, `StuckTracePTF`), executed together as a single `StatementSet` by `IsotopeReportsJob`. Sink topics use Apache Flink's `avro-confluent` format — SR-framed Avro, auto-registered on first write — so Control Center renders the rows natively. The application appears in CMF's applications API and Control Center's Flink tab. Drop everything (application + artifact + sink topics) with `make cp-flink-reports-down`.

Optional — fan-in provenance. Off by default. To also run the merge collector and its merge-edge sideband:

```
make cp-flink-reports-up ENABLE_MERGE_PROVENANCE=true
```

That adds two topics (`orders.flink_batched`, `isotope_merge_edge_markers`) and two more `INSERT`s to the same `StatementSet`. The merged records carry a fresh trace, and the edge topic records which parent traces fed each one — one row per contributing trace per window, weighted by `contributing_records`, a material fraction of the merge stage's write volume. See [docs/flink-collector.md §2.4](#). Teardown removes those topics either way.

[OPTIONAL] State-level provenance (§3.6). A second, independent switch:

```
make cp-flink-reports-up ENABLE_STATE_PROVENANCE=true

# both collectors – they answer different questions
make cp-flink-reports-up ENABLE_MERGE_PROVENANCE=true
ENABLE_STATE_PROVENANCE=true
```

That adds one topic (`isotope_state_provenance`) and one `INSERT` to the same `StatementSet`. Each traced order becomes an entity whose three stage records are three versions of it, chained by a `parents` array — `orders.placed` → `orders.enriched` → `orders.fulfilled`. Unlike the merge collector this

needs no window: identity is content-addressed, so it works on the updating sources (unbounded **GROUP BY**, joins, dedup) where the windowed path has nothing to key on. **CP only** — see [docs/state-provenance.md §5.0](#) for why CCAF cannot run this statement verbatim. Teardown removes the topic either way.

Watch it fill. The topic is **avro-confluent**, and **kafka-0** ships no Avro-aware consumer — **kafka-avro-console-consumer** lives in the Schema Registry pod, so run it there:

```
kubectll exec -n confluent schemaregistry-0 -- kafka-avro-console-consumer \
  --bootstrap-server kafka.confluent.svc.cluster.local:9071 \
  --property schema.registry.url=http://localhost:8081 \
  --topic isotope_state_provenance --from-beginning --max-messages 5
```

Both optional collectors have a guided walkthrough — [§7.1](#).

6.0 Drive traffic (required to see report rows)

All report jobs aggregate over **TUMBLE(event_time, INTERVAL '1' MINUTE)** windows, which only emit when the watermark advances past **window_end**. Spread records across **multiple** windows.

Start the three pipeline stages, each in its own terminal. They are long-running consumers that idle until traffic arrives:

```
./gradlew :app:run --args="enrich"           # orders.placed →
orders.enriched
./gradlew :app:run --args="fulfill"          # orders.enriched →
orders.fulfilled
./gradlew :app:run --args="ship"             # terminal consume
orders.fulfilled
```

Then, in a fourth terminal, drive the traffic through them:

```
# 30 records, 5s apart ≈ 2.5 min of event-time → spans 3+ windows
for i in {1..30}; do ./gradlew :app:run --args="place burst-$i" -q; sleep
5; done
```

Wait ~90s after the **last** record before checking results — that's the watermark catching up. **stuck_trace_alerts_1m** only fires for a trace that goes ≥60s of event time without a fresh hop; to exercise it, **stop the enrich/fulfill stages**, send one record to **orders.placed**, and keep sending unrelated records so the watermark crosses **event_time + 60s**.

7.0 Observe

```

make c3-open          # Control Center – report topics render natively
(Avro+SR)
make flink-ui         # Flink job graph / watermarks for the reports
(localhost:8082)
make metrics-up      # optional: Prometheus + Grafana for the stateless
meters

```

The metrics showcase has its own runbook — [k8s/monitoring/README.md](#) (and the meter/PromQL reference in [docs/metrics.md](#)). Stop the background port-forwards with `make c3-stop` / `make flink-ui-stop` / `make metrics-down`.

Ad-hoc Flink SQL — needs a session cluster (optional). The reports run as a CMF *Application* (§5.0): one compiled job under `execution.target=kubernetes-application`, which **cannot accept ad-hoc SQL**. Interactive querying needs a separate **session cluster**, and neither `cp-flink-up` nor `cp-flink-reports-up` creates one. Deploy it first:

```

make flink-deploy      # one-time: deploy the 'flink-basic' session
cluster (~40s)
make flink-sql        # interactive SQL Client on it, report DDL
preloaded

```

`flink-deploy` applies [k8s/base/flink-cluster-deployment.yaml](#) — a `cp-flink` 2.1 session cluster with 8 task slots, whose init containers fetch the Kafka and Avro-SR connector JARs. `flink-sql` then installs the PTF JAR into the JobManager's `/opt/flink/lib` and preloads the DDL from [scripts/flink/sql/cp/](#), so the session opens ready to query:

```

Flink SQL> SHOW TABLES;          -- 4 source tables + 7 report sinks +
views
Flink SQL> SHOW USER FUNCTIONS;  -- latency_percentiles, stuck_trace_ptf

```

Skip `flink-deploy` and `make flink-sql` stops with a message telling you to run it — it will not fall back to the reports Application pod. `make flink-ui` opens the reports Application by default; point it at this session cluster with `make flink-ui FLINK_UI_TARGET=flink-basic` to watch ad-hoc queries.

The SQL Client catalog is its own. The preloaded DDL lives in that session only, and the running reports Application registers its tables inside the job's own `TableEnvironment` — no shared catalog exists, so `SHOW TABLES` never reflects what the Application is running. Only DDL is preloaded; the report `INSERT INTO` files (`10_...70_`) are deliberately excluded, because the SQL Client rejects DML in an init file. Run them by hand only if you mean to — the Application is already writing to those same sink topics, so you would get duplicate rows, not an error.

7.1 [OPTIONAL] Showcase the two provenance collectors

Both collectors from §5.0 have to be on, and §6.0 traffic has to have run — including the `enrich` and `fulfill` stages, or every state entity has exactly one version and there is no chain to show:

```
make cp-flink-reports-up ENABLE_MERGE_PROVENANCE=true
ENABLE_STATE_PROVENANCE=true
```

Note about rendering the topics in Control Center: make `c3-open`, then Topics → any of `orders.flink_batched`, `isotope_merge_edge_markers`, `isotope_state_provenance` — all three are SR-framed Avro and render natively. If C3 will not come up, it is almost certainly the startup race in [KNOWN_ISSUES.md §1.0](#).

Showcasing the **two** provenance collectors takes both consumers, and they live in **different pods** — set them up before you present:

Topic	Format	Consumer	Pod
<code>orders.*</code> (sources, <code>flink_enriched</code>)	raw	kafka-console-consumer	kafka-0
<code>orders.flink_batched</code> , <code>isotope_merge_edge_markers</code> , <code>isotope_state_provenance</code>	avro-confluent	kafka-avro-console-consumer	schemaregistry-0

Both write a deprecation warning and a `--timeout-ms` shutdown trace to **stdout**, so both helpers keep only the data lines:

```
# Avro+SR sinks — schemaregistry-0 is the only pod with the Avro consumer.
avro() { kubectl exec -n confluent schemaregistry-0 -- kafka-avro-console-
consumer \
  --bootstrap-server kafka.confluent.svc.cluster.local:9071 \
  --property schema.registry.url=http://localhost:8081 \
  --from-beginning --timeout-ms 20000 "$@" 2>/dev/null | grep -E '^(\\{x-
isotope|NO_HEADERS)'; }

# raw-format order topics — kafka-0 is the only pod with the plain
consumer.
# grep -a matters: the values are binary, so without it grep reports
# "binary file matches" instead of printing the header line.
plain() { kubectl exec -n confluent kafka-0 -- kafka-console-consumer \
  --bootstrap-server localhost:9071 --from-beginning --timeout-ms 20000 \
  --property print.headers=true \
  --property
headers.deserializer=org.apache.kafka.common.serialization.StringDeserializ
er \
  "$@" 2>/dev/null | grep -a '^x-isotope'; }
```

7.1.1 Fan-in (Merge) Provenance — Three things to point at

7.1.1.1 The merged record carries a *fresh* trace, not a stolen one

```
avro --topic orders.flink_batched --max-messages 1 --property
print.headers=true \
  --property
headers.deserializer=org.apache.kafka.common.serialization.StringDeserializ
er \
| tr ',' '\n' | grep -aE 'trace-id|hop-count|origin-service|origin-ts'
```

```
x-isotope-trace-id:01a073036f007244ae76b286abd25c7f
x-isotope-origin-service:flink-batch
x-isotope-hop-count:1
x-isotope-origin-ts:1788636000000
```

The payload it is stamped on is the batch summary itself:

```
avro --topic orders.flink_batched --max-messages 1
```

```
{"pipeline":{"string":"orders"},"window_start":
{"long":1788635940000},"window_end":{"long":1788636000000},"event_count":
{"long":2},"distinct_traces":{"long":2}}
```

Now put the always-on 1:1 collector beside it:

```
plain --topic orders.flink_enriched --max-messages 1 \
| tr ',' '\n' | grep -aE 'trace-id|hop-count|origin-service|this-service'
```

```
x-isotope-trace-id:01a07302885b726faa1184fc79d6f984
x-isotope-origin-service:order-intake-service
x-isotope-this-service:flink-enrich
x-isotope-hop-count:2
```

That single digit is the whole argument. Same UDF family, same header shape, two different answers: the 1:1 forward **continued** an identity that **order-intake-service** minted (**hop_count** = 2, origin unchanged), and the merge **started** one (**hop_count** = 1, **origin_service** = **flink-batch**). Nobody told either statement which it was — the count says it. That is the [flink-collector.md §3.1](#) rule made visible: a trace is truthful only while every step has exactly one parent, so the merge does not get to claim one.

7.1.1.2 The merge trace ID is *derived*, not minted — the window is readable inside it

It is already visible in the output above: the trace ID opens with `01a073036f00`, and `0x01a073036f00` is `1788636000000` — the `window_end`. The first 12 hex characters are the UUIDv7 timestamp field, and the merge collector fills it from the window rather than from a clock:

```
01a073036f00 → 1788636000000 = window_end
01a073045960 → 1788636060000 = window_end
01a0730543c0 → 1788636120000 = window_end
```

This is why `80_merge_collector.fql` and `81_merge_edge_markers.fql` can agree on an ID without ever exchanging one, and why a replayed window reproduces it byte for byte.

7.1.1.3 Both reconciliations hold exactly

Join the edges to the merged records on `merge_trace_id` and check the two identities `flink-collector.md §2.4` asserts:

```
avro --topic orders.flink_batched --property print.headers=true \
  --property
headers.deserializer=org.apache.kafka.common.serialization.StringDeserializ
er \
| sed -E 's/.*x-isotope-trace-id:([0-9a-f]{32}).*\t(\{.*\})$/\1 \2/' >
/tmp/batched.txt
avro --topic isotope_merge_edge_markers > /tmp/edges.json

while read -r tid json; do
  ec=$(jq -r '.event_count.long' <<<"$json")
  dt=$(jq -r '.distinct_traces.long' <<<"$json")
  edges=$(grep -c "\"$tid\"" /tmp/edges.json)
  recs=$(grep "\"$tid\"" /tmp/edges.json | jq -s
'map(.contributing_records.long)|add')
  printf '%s edges=%-3s traces=%-3s records=%-3s events=%-3s %s\n' \
    "$tid" "$edges" "$dt" "$recs" "$ec" \
    "$([ \"$edges\" = \"$dt\" ] && [ \"$recs\" = \"$ec\" ] && echo OK || echo
MISMATCH)"
done < /tmp/batched.txt
```

```
01a073036f007244ae76b286abd25c7f edges=2   traces=2   records=2
events=2   OK
01a073045960714db1a35a4c0b28f701 edges=9   traces=9   records=9
events=9   OK
01a0730543c07df7b8b84e482fbd9c07 edges=10  traces=10  records=10
events=10  OK
```

`COUNT(*) == distinct_traces` and `SUM(contributing_records) == event_count`. Both hold because the two statements sit behind the same window operator — a late record is dropped by both or by neither.

What this demo cannot show. Every trace touches `orders.placed` exactly once, so `contributing_records` is always 1 and `event_count` always equals `distinct_traces`. The weighted-edge column is exercised in shape, not in value; a source where one trace contributes several records to a window is what makes the two reconciliations differ.

7.1.2 State Provenance — Three things to point at

7.1.2.1 The version chain, inline

One entity, three versions, each naming its predecessor — no join, no second topic:

```
avro --topic isotope_state_provenance > /tmp/state.json
E=$(head -1 /tmp/state.json | jq -r .entity_key.string)
jq -r --arg e "$E" 'select(.entity_key.string==$e)
  | "\(.source_name.string)\t\(.version_id.string)\tparents=\
(.parents.array|map(.string)|join(","))"' /tmp/state.json \
| column -t -s$'\t'
```

```
orders.placed      01a0730288817a44bb84484659ccfa6f  parents=
orders.enriched    01a07302bdde75c784f31084f187f5ea
parents=01a0730288817a44bb84484659ccfa6f
orders.fulfilled   01a07303085d7a6c8306afcc99abfb66
parents=01a07302bdde75c784f31084f187f5ea
```

The `DemoEvent` payload is forwarded verbatim across all three hops, so the **bytes never change** — and these are still three distinct versions, because `source_name` is part of the preimage ([state-provenance.md §2.1](#)).

7.1.2.2 The whole topic is internally consistent

Every parent resolves to a version of the same entity, one root per entity, nothing overflowed:

```
jq -s 'INDEX(.version_id.string) as $v
  | {versions: length,
    entities: (map(.entity_key.string)|unique|length),
    roots:    (map(select(.parents.array|length==0))|length),
    links:    (map(select(.parents.array|length>0))|length),
    every_parent_same_entity:
      (map(select(.parents.array|length>0)
        | .entity_key.string ==
$$.parents.array[0].string).entity_key.string) | all),
    overflow: (map(select(.parent_overflow.int>0))|length)}'
/tmp/state.json
```

```
{
  "versions": 93,
  "entities": 31,
  "roots": 31,
  "links": 62,
  "every_parent_same_entity": true,
  "overflow": 0
}
```

31 traced orders × 3 stages. Point out that the parent set is a **column of the record it describes**, so **versions** and **links** cannot drift the way **event_count** and the merge edge list could — that drift is not expressible here.

7.1.2.3 No hops anywhere

Ask the same consumer to print headers, and the records have none at all:

```
avro --topic isotope_state_provenance --max-messages 3 --property
print.headers=true \
  --property
headers.deserializer=org.apache.kafka.common.serialization.StringDeserializ
er \
| cut -c1-60
```

```
NO_HEADERS {"version_id":{"string":"01a0730288817a44bb844846
NO_HEADERS {"version_id":{"string":"01a07302bdde75c784f31084
NO_HEADERS {"version_id":{"string":"01a07303085d7a6c8306afcc
```

Not "no hop count" — **no isotope at all**. A row being revised is not a record taking a hop, so this collector declines the question rather than answering it wrong. Message lineage stays with the isotope on the **orders.*** topics; state lineage is a parallel record keyed by version. The two coexist without either lying.

7.1.3 The punchline — one order, three registers

The strongest single moment of the demo. **entity_key** in the state topic **is** the isotope trace ID, which **is** **contributing_trace_id** in the merge edges, so the same order is addressable in all three lineage models at once.

This continues from the blocks above — it reuses **\$E** from the state walkthrough and **/tmp/edges.json** from fan-in beat 3, so run those two first:

```
echo "trace $E"
plain --topic orders.flink_enriched --max-messages 1 | grep -ao "x-isotope-
hop-count: [0-9]*"
```

```
grep "\"$E\"" /tmp/edges.json \
| jq -r '"parent of merge_trace_id=\(.merge_trace_id.string) window=[\
(.window_start.long),\(.window_end.long)) records=\
(.contributing_records.long)'"
grep "\"$E\"" /tmp/state.json | jq -r '"version \(.source_name.string) = \
(.version_id.string)'"
```

```
trace 01a07302885b726faa1184fc79d6f984
x-isotope-hop-count:2
parent of merge_trace_id=01a073036f007244ae76b286abd25c7f window=
[1788635940000,1788636000000) records=1
version orders.placed = 01a0730288817a44bb84484659ccfa6f
version orders.enriched = 01a07302bdde75c784f31084f187f5ea
version orders.fulfilled = 01a07303085d7a6c8306afcc99abfb66
```

One order, read three ways: **as a message** it is an itinerary carried in-band; **as a fan-in parent** it is an edge recorded out-of-band, because the merge could not carry it; **as an entity** it is a version chain that never mentions hops. Three provenance models, one identity, and **no change to the isotope wire format** for any of them.

8.0 Teardown

Pick the depth:

```
make cp-flink-reports-down # drop reports / views / functions only
make flink-delete          # drop just the 'flink-basic' ad-hoc SQL
session cluster
make metrics-delete        # remove the Prometheus/Grafana showcase (pods
+ namespace)
make cp-flink-down         # Flink cluster + CMF + operator + cert-manager
(includes flink-delete)
make cp-down               # CP + operator (keeps minikube running)
make cp-teardown           # everything + stop minikube
make nuke                  # cp-teardown + minikube-delete + uninstall-
prereqs (factory reset)
```

9.0 Minimal path (no Flink)

To just watch a trace propagate without the Flink SQL reports:

```
make cp-up
make kafka-pf-up
./gradlew :app:run --args="place hello" # then enrich / fulfill / ship
```

Flink (§3.0 Flink) and the reports (§5.0 Deploy the 7 Flink reports (as a CMF Application)) are only needed for the SQL report topics.

10.0 Troubleshooting

- **Pods stuck Pending.** minikube is under-resourced — raise `MINIKUBE_CPUS` / `MINIKUBE_MEM` and `make minikube-delete && make cp-up`.
- **Flink job submission fails (services forbidden).** The supplemental RBAC didn't apply — `make flink-rbac`.
- **No report rows.** Almost always the watermark — see §6.0 Drive traffic (required to see report rows); traffic must span multiple 1-minute windows and you must wait ~90s after the last record.
- **App can't reach Kafka.** `make kafka-pf-up` isn't running, or the forward died — re-run it and confirm `localhost:30092` / `localhost:8081` are live.
- **make flink-sql says no session cluster / SHOW TABLES; returns an empty set.** Ad-hoc SQL needs the `flink-basic` session cluster — run `make flink-deploy` first, see §7.0 Observe. Confirm with `kubectl get flinkdeployment -n confluent: isotope-reports` alone means only the reports Application is up. Note that the report tables of the running Application are never visible to the SQL Client regardless — the catalogs are separate.
- **controlcenter-0 sits at 2/3 and make c3-open gets a connection refused.** C3 lost the startup race against Kafka's DNS and cannot recover on its own — cause and workaround in [KNOWN_ISSUES.md](#).
- **Control Center's Flink tab is blank.** CMF proxy connectivity — `make cmf-proxy-inject` (and `make cmf-proxy-logs` to debug).
- **make cp-flink-up times out waiting for the CMF pod / CMF CrashLoopBackOff.** Check `kubectl logs -n confluent -l app.kubernetes.io/name=confluent-manager-for-apache-flink`. If you see `Trial license ... expired / LicenseInitiator ... Constructor threw exception`, the image's embedded trial license has expired — supply your own via `CMF_LICENSE_SECRET` — see the CMF license note in §3.0 Flink. Wiping the CMF `PersistentVolumeClaim` does **not** reset it; the expiry is baked into the image.