

Confluent Kafka Isotope

`confluent-kafka-isotope` is a reference implementation of an **e-commerce order pipeline that uses Kafka Interceptors, Prometheus, and Apache Flink to capture, analyze, and report end-to-end event-tracing data in both batch and near real time.**

Much like isotopes used to trace molecules through a biochemical pathway, each event carries lightweight metadata that allows it to be followed as it moves through Kafka topics and distributed microservices.

Kafka topics become the connective tissue between services, while Kafka Interceptors quietly transform the event pipeline itself into an observable distributed system.

Table of Contents

- [1.0 How an Isotope Traverses an Event Pipeline](#)
 - [2.0 Project's Architecture](#)
 - [3.0 Project layout](#)
 - [4.0 Getting Started](#)
 - [4.1 Unit tests \(no broker, instant\)](#)
 - [4.2 Demo CLI — see one trace propagate live](#)
 - [4.2.1 Full Bipartite Demo](#)
 - [4.3 Integration tests \(live Kafka via Minikube\)](#)
 - [4.4 Confluent Platform on Minikube — Production-Like Streaming, Running Locally](#)
 - [4.5 Flink SQL reports on Confluent Cloud for Apache Flink \(CCAF\)](#)
 - [4.6 Stateless reports via Micrometer + Prometheus/Grafana \(optional\)](#)
 - [5.0 Resources](#)
-

1.0 How an Isotope Traverses an Event Pipeline

An **Isotope** is a **lightweight tracing artifact attached to Kafka record headers**. Like a biochemical isotope used to trace molecules through a metabolic pathway, it allows the journey of a record through an event-driven architecture to be observed and analyzed.

This project models a Kafka pipeline as a **bipartite graph**: *services occupy one vertex set, topics the other, and every produce and consume operation forms an edge between them*. The resulting graph provides a unified view of the complete event topology, from producers to intermediate processing services to terminal consumers.

A **ProducerInterceptor** stamps each record with an isotope and appends one hop for every `send()` operation, capturing the produce edges. Consumers call `IsotopeContext.recordConsume()` to emit a lightweight marker representing the corresponding consume edges, while services that consume and then produce invoke `IsotopeContext.adoptFromRecord()` so the trace identity persists across every hop. Apache Flink reconstructs the complete **service → topic → service** graph from the isotope headers alone, producing seven reports: end-to-end latency, latency percentiles, produce-side topology, the full bipartite topology, hop distribution, per-topic coverage (a trace-loss funnel signal), and stuck-trace

detection. The same implementation runs unchanged on both Confluent Platform (self-managed) with Confluent Manager for Apache Flink (CMF) and Confluent Cloud for Apache Flink (CCAF).

Full architectural design of Isotope Tracing — the header layout (`x-isotope` JSON + seven scalar headers with a worked example), how the producer interceptor gets invoked, why the consume side uses explicit calls instead of a `ConsumerInterceptor`, and the bipartite-graph rationale — is in [docs/design.md](#).


2.0 Project's Architecture

A bird's-eye view of the moving parts. The demo CLI in `app/` consumes the external tracing library (`ai.signalroom:kafka-isotope-core`), which registers a Kafka `ProducerInterceptor` that stamps an isotope into record headers on every `send()`. Consume-then-produce services propagate the inbound trace by explicitly calling `IsotopeContext.adoptFromRecord(record)`. Business events then flow through a three-topic Kafka pipeline, where Flink SQL reads the isotope metadata and emits one-minute aggregate reports.

Both runtimes run the same seven logical reports off the same source and view definitions, though each runtime keeps its own copy of the SQL. **Confluent Platform (CP)** on Minikube executes the `.fql` files under `scripts/flink/sql/cp/` as a single Flink 2.1 Confluent Manager for Apache Flink (CMF) Application (`IsotopeReportsJob`), while **Confluent Cloud for Apache Flink (CCAF)** applies the same logical SQL as inline `confluent_flink_statement` Terraform resources under `terraform/`. The shadow JAR from `ptf/`—which powers two of the seven reports—runs unchanged on both runtimes: bundled into the CP application JAR and uploaded as a Flink artifact on CCAF.

Alongside that—**additive, opt-in, and disabled by default**—the interceptor can also emit **Micrometer** metrics for **Prometheus**, with **Grafana** providing visualization (§4.6). This path produces the three **stateless** reports (`latency_1m`, `topology_1m`, and `hop_distribution_1m`) without requiring a stream processor. The remaining four reports continue to run in Flink because they depend on per-`trace_id` state or absence-of-event analysis, which falls outside Prometheus's query model.

(Kafka is drawn once below for brevity; each runtime provisions its own Kafka cluster.)



See [§3.0](#) for the file tree behind each box, and [§4.0](#) for the run commands.

3.0 Project layout

The isotope tracing library lives in its own repo — [j3-signalroom/kafka-isotope](#) ([ai.signalroom:kafka-isotope-core](#) + [ai.signalroom:kafka-isotope-metrics](#)). This repo is the runnable demo that consumes it.

app/ isotope library)	demo CLI + tests (consumes the
src/main/proto/ai/signalroom/kafka/isotope/proto/ demo_event.proto	DemoEvent message (Protobuf value
schema)	
src/main/java/ai/signalroom/kafka/isotope/ App.java	demo CLI – pipeline-position verbs
	(place / enrich / fulfill / ship) +
generic	send / hop / consume / sink modes;
registers	IsotopeProducerInterceptor + starts
the	PrometheusIsotopeMetrics exporter
(§4.6)	
src/integrationTest/java/.../ ProducerInterceptorIT,	BrokerSmokeIT,
BipartiteTopologyIT,	ThreeStageHopPropagationIT,

tests; produce/consume	IsotopeTestHarness – live-broker
forwarded)	DemoEvent via SR-framed Protobuf (need Minikube CP + SR port-
ptf/ shadow JAR	Flink reports application + PTF
src/main/java/ai/signalroom/kafka/isotope/flink/ IsotopeReportsJob.java	CP entry point – reads the bundled
sql/*.fql,	registers both PTFs
programmatically, runs the	7 INSERT INTOs as one StatementSet
(CMF Application)	T-Digest p50/p95/p99 (PTF: per-
LatencyPercentilesPTF.java	per-trace state + event-time timer
window state + timers)	shared T-Digest (de)serialization
StuckTracePTF.java	TDigestsTest
TDigests.java	CFK / CMF manifests (applied by
src/test/java/.../ k8s/base/ `make cp-up` / `flink-up`)	Kafka / SR / Connect / ksqlDB /
confluent-platform-c3++.yaml	in-cluster S3-compatible store
Control Center	cmf:// artifact (JAR) storage
minio.yaml	Helm values for CMF 2.4 – artifact
backing CMF's	(points at MinIO) + writable
cmf-values.yaml	FlinkApplication template
storage	deploy-cmf-flink-reports.sh) – the
environment catalog	custom cp-flink image for the CMF
cmf-flink-application.json	bakes in the Kafka + avro-confluent
(envsubst'd by	and the s3-fs-hadoop plugin (`make
reports job	optional cp-flink session cluster
flink-sql-isotope.Dockerfile	`make flink-deploy` / `make flink-
compute pool:	(not used by the reports path)
SQL connectors	RBAC for the cp-flink operator
flink-image-build`)	optional metrics showcase (§4.6) –
flink-cluster-deployment.yaml	dedicated 'monitoring' namespace
for ad-hoc	Prometheus pod/Service; scrapes
sql`	via
flink-rbac.yaml	Grafana pod/Service; auto-
k8s/monitoring/ `make metrics-up`	
00-namespace.yaml	
10-prometheus.yaml	
host stages	
host.minikube.internal:9410/9411/9412	
20-grafana.yaml	

provisioned datasource

produce/consume meters

kustomization.yaml

README.md

scripts/

port-forward-kafka.sh

localhost:8081 → SR

port-forward-taskmanager.sh

deploy-cmf-flink-reports.sh

cmf://

Flink 2.1

deploy-cc-flink-reports.sh

`terraform apply`

cc-cli-env.sh

`terraform output`,

BOOTSTRAP /

JAAS / ...

cc-app-run.sh

:app:run` that

the six -D flags

flink/README.md

(CP=7 reports/Avro+SR,

layout, operations

flink/sql/cp/

01_register_functions,

06_consume_events_view,

reports, 99_teardown

terraform/setup-confluent-flink.tf.)

flink/sql/ccaf-ai/

off by default)

trace_rca.fql

turns each

cause

ML_PREDICT

ai.tf)

terraform/

cc-flink-reports-up`)

providers.tf

key/secret vars

+ 8-panel dashboard for the 6

`kubectl apply -k k8s/monitoring`

runbook + troubleshooting

localhost:30092 → Kafka,

Flink TaskManager web UI forward

builds shadow JAR + uploads it as a

artifact + deploys the reports as a

CMF Application (runs sql/cp/*.fql)

builds shadow JAR + wraps

for the CCAF path

pulls Kafka + SR creds from

builds the JAAS string, exports

SR_URL / KAFKA_KEY / KAFKA_SECRET /

thin wrapper around `./gradlew

sources cc-cli-env.sh and injects

Flink SQL reports – runtime split

CCAF=7 reports/Protobuf+SR),

CP Flink SQL: 00_source_table,

05_isotope_view,

05_report_sinks (avro-confluent),

10/20/25/30/40/60/70 INSERT INTO

(CCAF SQL is inlined under

optional AI extension (CCAF-only,

PoC: 8th, AI-generated report –

stuck-trace alert into an LLM root-

hypothesis via CREATE MODEL +

(wired in by terraform/setup-ccaf-

CCAF infrastructure-as-code (`make

Confluent provider – cloud

```

versions.tf
provider versions
variables.tf
region, day_count
data.tf
sources
  setup-confluent-environment.tf
governance package)
  setup-confluent-kafka.tf
rotation module

tf_module)
  setup-confluent-flink.tf
compute pool,

rotation, and 25 inline

resources: 4 ALTER TABLE

DROP FUNCTION +

INSERT INTO
  setup-ccaf-ai.tf
extra

sink +

false)
  outputs.tf
rotating

(sensitive)
docs/
from the README)
  design.md
deep-dive)
  isotope_diagram.png
isotope and its flow through Kafka headers
  runbook-minikube.md
($4.4)
  runbook-ccaf.md
($4.5)
  metrics.md
PromQL reference ($4.6)
  terraform.png
in $4.5)
Makefile
flink-reports-up /

reports-down /

metrics-delete / ...

```

```

required Terraform (>= 1.13) +

confluent_api_key/secret, cloud,

organization lookup + other data

environment (ESSENTIALS stream-

Kafka cluster + Kafka API key

(iac-confluent-api_key-rotation-

service account + 6 role bindings,

artifact upload, SR API key

`confluent_flink_statement`

+ 3 VIEW + 7 sink CREATE TABLE + 2

2 CREATE FUNCTION (both PTFs) + 7

OPTIONAL AI trace-RCA report - 3

statements (CREATE MODEL + Protobuf

INSERT ... ML_PREDICT); gated on

var.enable_trace_rca (default

environment_id, bootstrap, SR URL,

Kafka + SR API key/secret outputs

extracted long-form docs (linked

how the isotope is carried ($1.0

visual representation of the

full CP-on-Minikube run sequence

full CCAF / Terraform run sequence

Micrometer/Prometheus meter +

rendered resource graph (embedded

cp-up / flink-up / kafka-pf-up /

cc-flink-reports-up / cc-flink-

metrics-up / metrics-down /

```

4.0 Getting Started

First, clone the repo to your local machine using the [GitHub CLI](#):

```
gh repo clone j3-signalroom/confluent-kafka-isotope
```

Change to the repo directory:

```
cd /path/to/confluent-kafka-isotope
```

Then decide what you want to run:

4.1 Unit tests (no broker, instant)

```
./gradlew test # runs :ptf:test (the app has no unit tests in this repo)
# or scoped:
./gradlew :ptf:test # TDigestsTest – T-Digest sketch
(de)serialization + accuracy
```

4.2 Demo CLI — see one trace propagate live

The fastest way to watch the isotope mechanic. Requires the cluster to be up and the Kafka + SR forwards running (see step 3 below for the bring-up commands). The CLI has two argument styles — **pipeline-position verbs** that bake in the orders.* topic chain (recommended for the demo) and **generic verbs** that take raw topic + service args (for ad-hoc inspection on any topic):

Verb	Args	What it does
place	[payload]	Produces one isotope-tagged DemoEvent to orders.placed as order-intake-service (default payload: hello), then exits. Auto-creates the topic.
enrich	—	Consumes from orders.placed , adopts the isotope, emits a consume-edge marker to isotope_consume_edge_markers as order-enrichment-service , then re-produces to orders.enriched . Runs until Ctrl-C.
fulfill	—	Same as enrich but for orders.enriched → orders.fulfilled as order-fulfillment-service . Runs until Ctrl-C.
ship	—	Terminal consumer for orders.fulfilled as shipping-notification-service . Emits a consume-edge marker so it

Verb	Args	What it does
		shows up in the bipartite report; does not re-produce. Runs until Ctrl-C.
send	<topic> <service> <payload>	Generic produce. Auto-creates the topic.
hop	<in-topic> <out-topic> <service>	Generic consume-then-produce; emits a consume-edge marker. Runs until Ctrl-C.
consume	<topic> <service>	Generic terminal-consume; emits a consume-edge marker and pretty-prints the trail. Runs until Ctrl-C.
sink	<topic>	Passive peek — pretty-prints the isotope trail but does NOT emit a consume marker. Use for ad-hoc inspection. Runs until Ctrl-C.

4.2.1 Full Bipartite Demo

A 4-stage chain (full bipartite graph) in four terminals. Run them in pipeline order — **place** produces, then **enrich** / **fulfill** / **ship** each pick up where the previous stage left off:

```
# Terminal A — kick the chain off (run repeatedly to send more)
./gradlew :app:run --args="place 'hello world'" -q

# Terminal B — first hop: orders.placed → orders.enriched
./gradlew :app:run --args="enrich" -q

# Terminal C — middle hop: orders.enriched → orders.fulfilled
./gradlew :app:run --args="fulfill" -q

# Terminal D — terminal consumer (prints the full 3-hop trail AND emits a
#                          consume-edge marker so shipping-notification-service shows
up
#                          in the bipartite report)
./gradlew :app:run --args="ship" -q
```

Terminal D's output for each record shows the same **trace_id** across all three hops, **origin = order-intake-service** (never reassigned), and **hops[]** listing **order-intake-service → order-enrichment-service → order-fulfillment-service** in order with per-hop timestamps. The bipartite-topology report sees all six edges: produce edges **order-intake-service → orders.placed**, **order-enrichment-service → orders.enriched**, **order-fulfillment-service → orders.fulfilled** and consume edges **orders.placed → order-enrichment-service**, **orders.enriched → order-fulfillment-service**, **orders.fulfilled → shipping-notification-service**. Swap **consume** for **sink** if you only want to inspect records without recording the terminal edge. Override endpoints via **-Dkafka.bootstrap=...** / **-Dschema.registry.url=...** if you're not on the default Minikube layout.

Collector

Stamp headers; emit markers + metrics



Four-stage order workflow; traces ride in headers.
shipping-notify can't append a header (no re-produce),
so it emits a consume-edge marker to its own topic.
The topology graph stays complete either way.

Interpreter

Reads headers + markers



7 reports · CP + CCAF
percentiles is a PTF

- | | | |
|--|--|---------------------------------------|
| ☐ Service | ☐ Event topic (x-isotope) | ☐ Marker topic |
| ☐ Flink interpreter | ☐ Report sink (1-min) | ☐ Shadow JAR |
| ☐ Metrics backend | — Produce | - - - Edge marker emit |

Just want the commands? The full local stack-up sequence (cluster → Kafka → Flink → reports → traffic → teardown) is consolidated in [docs/runbook-minikube.md](#).

4.3 Integration tests (live Kafka via Minikube)

Bring up the local Confluent Platform stack and port-forward Kafka + SR:

```

make minikube-start          # one-time
make cp-up                   # CFK Operator +
Kafka/SR/Connect/ksqlDB/C3 (~5 min)
make kafka-pf-up             # localhost:30092 → Kafka,
localhost:8081 → Schema Registry
  
```

Then run the suite:

```

./gradlew :app:integrationTest          #
every IT below
./gradlew :app:integrationTest --tests '*ProducerInterceptorIT'  #
just one
  
```

Override the endpoints if needed:

```
./gradlew :app:integrationTest \
  -PkafkaBootstrap=localhost:30092 \
  -PschemaRegistryUrl=http://localhost:8081
```

Tear down forwards when done:

```
make kafka-pf-down
```

The integration tests cover:

Test	What it verifies
BrokerSmokeIT	AdminClient can create/list/delete a topic via the NodePort port-forward
ProducerInterceptorIT	A consumer sees the <code>x-isotope</code> JSON header + all 7 scalar reporting headers with the expected origin/hop values, and the Protobuf round-trip preserves <code>DemoEvent.source</code> / <code>payload</code>
ThreeStageHopPropagationIT	<code>order-intake-service</code> → <code>topic-AB</code> → <code>order-enrichment-service</code> → <code>topic-BC</code> → <code>order-fulfillment-service</code> produces a stable trace ID, 2-hop trail in send order, and correct scalar headers (origin = <code>order-intake-service</code> , this = <code>order-enrichment-service</code> , hop count = 2) at the terminal; consume-then-produce hops use <code>IsotopeContext.adoptFromRecord</code> to carry the trace forward
BipartiteTopologyIT	The 4-stage <code>order-intake-service</code> → <code>topic-AB</code> → <code>order-enrichment-service</code> → <code>topic-BC</code> → <code>order-fulfillment-service</code> → <code>topic-CD</code> → <code>shipping-notification-service</code> chain emits exactly three consume-edge markers to a per-test markers topic — one per consume edge. Every marker carries the trace ID, forwarded <code>x-isotope-*</code> scalars describing the upstream producer, and the new <code>x-isotope-consumer-service</code> naming the downstream consumer. Asserts the <code>(consumer_service, consumed_topic)</code> set is exactly the three pairs of stages 2-4

4.4 Confluent Platform on Minikube — Production-Like Streaming, Running Locally

Caveat: Minikube is a single-node cluster. It is not a production-like environment, but it is sufficient for local development and testing. The Confluent Platform (CP) stack runs in Minikube, and you can port-forward Kafka and Schema Registry to your local machine.

To **run**, **test**, and **debug** Apache Flink like a production engineer, this project provides a full Confluent Platform stack running locally on [Minikube](#) — no cloud required.

You get a production-like environment on your machine, with all the components you'd expect in a real deployment:

- **Confluent Platform** (KRaft mode) via Confluent for Kubernetes (CFK)
- **Apache Flink 2.1.2** via the Confluent Flink Kubernetes Operator 1.140.1
- **Confluent Manager for Apache Flink (CMF) 2.4.0** for Flink environment management

To run this project, you'll need **macOS (with Homebrew)** or **Linux (with apt-get)**. The full stack — **Minikube + Confluent Platform + Flink + CMF** — is resource-intensive and designed to mirror a production environment. Therefore, the following defaults are recommended:

Resource	Default
CPU	6
Memory	20 GB
Disk	50 GB

These settings ensure stable performance across all components. You can tune them as needed, but lower resource levels may cause pod restarts or degraded performance.


Seven reports — five pure Flink SQL plus two JAR-backed PTFs — run as a **single Flink 2.1 CMF Application** ([IsotopeReportsJob](#), one StatementSet, seven sinks) submitted through CMF and executed by the Confluent Flink Kubernetes Operator. No raw session cluster is involved; [k8s/base/flink-cluster-deployment.yaml](#) (`make flink-deploy` / `make flink-sql`) is a separate, optional path for ad-hoc SQL. Confluent Cloud runs the same seven reports from its own copy of the SQL, inlined in Terraform — see [§4.5](#) for that path; this section is the local-Minikube one.

The full bring-up sequence — cluster → Flink → reports → traffic → teardown — is consolidated in [docs/runbook-minikube.md](#). The short version: `make flink-up` then `make flink-reports-up`, then drive traffic across **multiple** 1-minute windows (a single burst sits in one open window forever — the watermark has to cross `window_end` for a tumbling window to emit) and wait ~90s after the last record.

Report sink topics ride **Avro+SR** ([avro-confluent](#), auto-registered on first write) so Control Center renders them natively — a deliberate *format-by-domain* split: app events are **Protobuf+SR** ([DemoEvent](#)), Flink aggregates are **Avro+SR** (cp-flink ships no SR-integrated Protobuf format), and the consume-edge marker topic [isotope_consume_edge_markers](#) is **null-value / headers-only**. Not a defect — a clean split by domain.

4.5 Flink SQL reports on Confluent Cloud for Apache Flink (CCAF)

The CCAF parallel of [§4.4](#), driven by Terraform under [terraform/](#) — no local cluster. One `make` target spins up a fresh Confluent Cloud environment, Kafka cluster, compute pool, two rotating service-account API key pairs (Kafka + Schema Registry), the uploaded PTF JAR, and 25 `confluent_flink_statement` resources (4 ALTER TABLE + 3 VIEW + 7 sink CREATE TABLE + 2 CREATE FUNCTION + 7 INSERT INTO, plus 2 transient DROP FUNCTION). The shape mirrors [apache_flink-kickstarter-ii](#) — same provider version and `iac-confluent-api_key_rotation-tf_module`.



The full provision → deploy → traffic → teardown sequence is consolidated in [docs/runbook-ccaf.md](#). The short version: **export** a Cloud-level Confluent Cloud API key, **make cc-flink-reports-up** (~6–8 min first run; idempotent re-applies), drive traffic with **scripts/cc-app-run.sh place|enrich|fulfill|ship** across **multiple** 1-minute windows (a single burst sits in one open window forever, and you wait ~90s after the last record), then **make cc-flink-reports-down** to **terraform destroy** the whole environment.

Two CCAF-specific design points worth keeping in the README:

Format-by-runtime (not -by-domain). CP's reports land on **Avro+SR** (`'value.format' = 'avro-confluent'` in [scripts/flink/sql/cp/05_report_sinks.fql](#)). CCAF's reports land on **Protobuf+SR** (`'value.format' = 'proto-registry'` in each sink's WITH clause in [terraform/setup-confluent-flink.tf](#)). The two runtimes' SQL is otherwise unshared: CP's lives hardcoded in [scripts/flink/sql/cp/](#), CCAF's lives inline as `confluent_flink_statement` resources in [terraform/setup-confluent-flink.tf](#).

Why percentiles is a PTF. CCAF rejects all **CREATE FUNCTION** statements for user-defined *aggregate* functions ("aggregate functions are not supported"). Percentiles would naturally be an aggregate, so to keep the report portable it's implemented as a `ProcessTableFunction` instead — `LATENCY_PERCENTILES` (class `LatencyPercentilesPTF`) does its own 1-minute tumbling-window aggregation over a T-Digest sketch via per-window state and event-time timers. A PTF has no such restriction, so it registers and runs on **both** runtimes, exactly like `STUCK_TRACE_PTF`. Both runtimes therefore run the same seven reports: **latency** (avg/min/max), **topology** (produce-side), **bipartite_topology** (full service↔topic↔service graph), **hop_distribution**, **coverage**, **stuck_trace**, and **latency_percentiles** (p50/p95/p99).

Optional: AI root-cause analysis (CCAF-only, off by default). Beyond the seven deterministic reports, [terraform/setup-ccaf-ai.tf](#) wires an **eighth, AI-generated report** that turns each stuck-trace *alert* into a natural-language root-cause hypothesis plus a one-line remediation. It adds three `confluent_flink_statement` resources — `CREATE MODEL trace_rca` (a remote text-generation model), a `proto-registry` sink `isotope_report_trace_rca_1m`, and an `INSERT ... SELECT ... LATERAL TABLE(ML_PREDICT('trace_rca', ...))` — all **gated on `var.enable_trace_rca` (default `false`)**, so a normal **make cc-flink-reports-up** is completely unaffected. The model is invoked once per *alert* (the low-volume 1-minute stuck-trace report topic), never per record, and its output lands on its own topic — it never overwrites the deterministic reports. Defaults target OpenAI (`gpt-4o`); for Claude hosted by Anthropic — a **direct** CCAF provider — set `rca_model_provider = "anthropic"`, `rca_model_endpoint = "https://api.anthropic.com/v1/messages"`, a Claude `rca_model_version`, and your Claude `rca_model_api_key` (bare API key, no AWS auth), plus the Anthropic-required `rca_model_max_tokens`. Claude via AWS Bedrock (`rca_model_provider = "bedrock"`) also works but isn't required. The standalone SQL PoC — a `CREATE MODEL + ML_PREDICT` walkthrough with provider notes and alternative options — is in [scripts/flink/sql/ccaf-ai/trace_rca.fql](#).

4.6 Stateless reports via Micrometer → Prometheus/Grafana (optional)

Three of the seven reports — `latency_1m`, `topology_1m`, `hop_distribution_1m` — are pure stateless scalar aggregation over bounded-cardinality dimensions (service / topic / hop_count, never `trace_id`). Those don't need a stream processor: the `producer interceptor` already has every value in scope on each `send()`, so it emits them as **Micrometer** meters (`PrometheusIsotopeMetrics.java`) and lets **Prometheus** window them at read time, with **Grafana** on top. The other four (`latency_percentiles`, `coverage`, `bipartite_topology`, `stuck_trace`) are per-`trace_id` stateful or absence-of-event problems Prometheus can't express, so they **stay in Flink**. This path is **additive and opt-in** — a 3-Micrometer / 4-Flink split.

Full details — the meter/PromQL reference, the produce- and consume-side signals, the two deliberate gaps (`distinct_traces`, windowed `min`), and why percentiles stays a PTF — are in [docs/metrics.md](#). The one-command Prometheus + Grafana showcase has its own runbook: [k8s/monitoring/README.md](#) (`make metrics-up`).

5.0 Resources

- [Medium Article: Kafka's quiet observability superpower — Kafka Interceptors](#)