

Flink SQL reports

Seven deterministic reports (plus an optional eighth, AI-generated one — see [§2.0 \[OPTIONAL\] AI Root-Cause Analysis \(RCA\)](#)) that read two streams of isotope metadata — the **produce side** (the `x-isotope-*` headers stamped on every event topic record by `IsotopeProducerInterceptor`) and the **consume side** (the value-less marker records on `isotope_consume_edge_markers` emitted by `IsotopeContext.recordConsume`) — and surface what's flowing where, how fast, and how reliably. Each report runs as a long-lived `INSERT INTO <report>_1m SELECT ...` streaming job. The aggregation logic is identical across runtimes; only the source/sink DDL and function-registration glue differ.

Riding alongside them is one **collector** `INSERT` — not a report, and not an aggregation. Where the seven reports *interpret* isotope headers that an interceptor already wrote, the collector *writes* them: it forwards `orders.placed` to `orders.flink_enriched` 1:1, appending a Flink hop so Flink appears in the topology graph as a producer rather than only a reader. Both runtimes run all eight `INSERT`s as a single job. See [docs/flink-collector.md](#).

The headline addition is the `bipartite_topology` report: it unions the produce and consume views to render the pipeline as a literal **bipartite graph** from graph theory — services in one vertex set, topics in the other, edges crossing between them in both directions. See [root README §1.0 "How an Isotope Traverses an Event Pipeline"](#) for the full motivation.

Runtime	Reports	Sink format	Where the SQL lives
Confluent Platform Flink (Flink 2.1 CMF Application on Minikube)	7 (latency, topology, bipartite-topology, hop-distribution, coverage, stuck-trace, latency-percentiles) + 1 collector	<code>avro-confluent</code> (SR-framed Avro)	<code>scripts/flink/sql/cp/</code> — run by <code>IsotopeReportsJob</code> as one <code>StatementSet</code> , deployed by <code>scripts/deploy-cmf-flink-reports.sh</code>
Confluent Cloud for Apache Flink (CCAF)	7 (same set as CP) + 1 collector	<code>proto-registry</code> (SR-framed Protobuf)	Inlined as <code>confluent_flink_statement</code> resources in <code>terraform/setup-confluent-flink.tf</code> , the <code>INSERT</code> s consolidated into one <code>EXECUTE STATEMENT SET</code> — applied by <code>scripts/deploy-cc-flink-reports.sh</code>

Control Center deserializes both sink formats natively.

Table of Contents

- [1.0 What each report computes](#)
- [2.0 \[OPTIONAL\] AI Root-Cause Analysis \(RCA\)](#)
- [3.0 Layout](#)
- [4.0 Wire-format detail \(CP only\)](#)

- [5.0 PTF JAR](#)
- [6.0 Operations](#)
 - [6.1 CP Flink on Minikube](#)
 - [6.2 CCAF](#)

1.0 What each report computes

Every report aggregates over a `TUMBLE(event_time, INTERVAL '1' MINUTE)` window. The "Source view" column names the typed view (or views) the report reads from; the views in turn filter the single `isotope_raw` source table by the presence/absence of the `x-isotope-consumer-service` header.

Every report also carries a `pipeline` column (decoded from the `x-isotope-pipeline` header) as a leading grouping dimension, so rows for an `orders` pipeline and a `location` pipeline never aggregate together. The "What it computes" column below lists the *additional* keys each report groups by, on top of `pipeline`.

Report	Source view	What it computes	Runtimes
<code>latency</code>	<code>isotope</code>	avg / min / max end-to-end latency by <code>origin_service</code> × <code>current_topic</code>	CP, CCAF
<code>topology</code>	<code>isotope</code>	produce-edge counts per (<code>producer_service</code> , <code>topic</code>) per minute	CP, CCAF
<code>bipartite_topology</code>	<code>isotope</code> and <code>consume_events</code>	full bipartite graph: produce edges plus consume edges per minute	CP, CCAF
<code>hop_distribution</code>	<code>isotope</code>	record counts bucketed by <code>hop_count</code> per topic per minute	CP, CCAF
<code>coverage</code>	<code>isotope</code>	distinct traces per topic per minute	CP, CCAF
<code>stuck_trace</code>	<code>isotope</code>	alerts (via <code>STUCK_TRACE_PTF</code>) for traces idle ≥60s of event time	CP, CCAF
<code>latency_percentiles</code>	<code>isotope</code>	p50 / p95 / p99 (via <code>LATENCY_PERCENTILES</code> PTF, T-Digest)	CP, CCAF

Plus one non-report INSERT. `75_flink_collector.fql` is not an aggregation and has no window — it forwards `orders.placed` to `orders.flink_enriched` 1:1, appending a Flink hop via the `ISOTOPE_APPEND_HOP` UDF. It rides in the same job as the seven reports on both runtimes. See [docs/flink-collector.md](#).

Format-by-runtime (not-by-domain). CP's reports land on **Avro+SR** (`'value.format' = 'avro-confluent'` in [scripts/flink/sql/cp/05_report_sinks.fql](#)). CCAF's reports land on **Protobuf+SR** (`'value.format' = 'proto-registry'` in each sink's `WITH` clause in [terraform/setup-confluent-](#)

[flink.tf](#)). The two runtimes' SQL is otherwise unshared: CP's lives hardcoded in [scripts/flink/sql/cp/](#), CCAF's lives inline as `confluent_flink_statement` resources in [terraform/setup-confluent-flink.tf](#).

Why `latency_percentiles` is a `ProcessTableFunction` (PTF). Since CCAF does not support `User-Defined AGGregate functions` (UDAGG), I implemented the percentiles report as a `ProcessTableFunction` to keep it portable across runtimes. `LATENCY_PERCENTILES` (class `LatencyPercentilesPTF`) performs its own 1-minute tumbling-window aggregation over a T-Digest sketch using per-window state and event-time timers. **A PTF avoids that UDAGG restriction, so it registers and runs on both runtimes, just like `STUCK_TRACE_PTF`.** Both runtimes therefore run the same seven reports: `latency` (avg/min/max), `topology` (produce-side), `bipartite_topology` (full service↔topic↔service graph), `hop_distribution`, `coverage`, `stuck_trace`, and `latency_percentiles` (p50/p95/p99).

The demo event topics (`orders.placed`, `orders.enriched`, `orders.fulfilled`) ride Protobuf+SR via the Java app's `DemoEvent` schema on both runtimes — those are written by the Kafka producer client, not by Flink, so the SR-Protobuf gap doesn't apply.

2.0 [OPTIONAL] AI Root-Cause Analysis (RCA)

Beyond the seven deterministic reports, an **eighth, AI-generated report** turns each stuck-trace *alert* into a natural-language root-cause hypothesis plus a one-line remediation. This **CCAF-only** feature is wired by [terraform/setup-ccaf-ai.tf](#) and can be enabled **at deploy time** by setting `ENABLE_TRACE_RCA=true` and supplying the other required `RCA` arguments to the `make cc-flink-reports-up` target.

When enabled, it adds three `confluent_flink_statement` resources:

- `CREATE MODEL trace_rca` — registers a remote text-generation model.
- A `proto-registry` sink table `isotope_report_trace_rca_1m`.
- An `INSERT ... SELECT ... LATERAL TABLE(ML_PREDICT('trace_rca', ...))` that calls the model **once per alert** (reading the low-volume 1-minute stuck-trace report topic, never the source stream) and writes the hypothesis to its own topic — it never overwrites the deterministic reports.

The default provider is OpenAI (`gpt-4o`). Claude is supported via **AWS Bedrock** (`rca_model_provider = "bedrock"` with AWS credentials). Other supported providers: `vertexai`, `azureopenai`, `googleai`, `sagemaker`, `azureml`.

3.0 Layout

<code>scripts/flink/sql/cp/</code>	CP Flink – session-cluster SQL
<code>00_source_table.fql</code>	CREATE TABLE per topic ('connector'
<code>= 'kafka') + isotope_raw UNION view</code>	
<code>01_register_functions.fql</code>	CREATE FUNCTION ... USING JAR
<code>'file:///opt/flink/lib/isotope-flink-udf.jar'</code>	
<code>05_isotope_view.fql</code>	Typed view; decodes x-isotope-*
<code>header scalars (produces only)</code>	
<code>06_consume_events_view.fql</code>	Typed view of
<code>isotope_consume_edge_markers markers (consume edges)</code>	
<code>05_report_sinks.fql</code>	CREATE TABLE for each

```

isotope_report_*_1m Kafka sink (avro-confluent)
  10_latency_report.fql          INSERT INTO: avg/min/max latency by
origin × topic
  20_topology_report.fql        INSERT INTO: produce-edge counts
per minute
  25_bipartite_topology_report.fql INSERT INTO: produce plus consume
edges (bipartite graph) per minute
  30_hop_distribution.fql        INSERT INTO: hop-count buckets per
topic per minute
  40_coverage_report.fql        INSERT INTO: distinct traces per
topic per minute
  60_stuck_trace_report.fql      INSERT INTO: stuck-trace alerts via
STUCK_TRACE_PTF
  70_latency_percentiles_report.fql INSERT INTO: p50/p95/p99 via
LATENCY_PERCENTILES PTF (T-Digest)
  07_flink_collector_sink.fql    CREATE TABLE orders.flink_enriched
– headers column declared

makes it writable
  75_flink_collector.fql        INSERT INTO: Flink stamps its own
hop via ISOTOPE_APPEND_HOP

window)
  08_merge_provenance_sinks.fql  OPTIONAL – CREATE TABLE
orders.flink_batched (writable headers)
+ isotope_merge_edge_markers (typed
edge list)
  80_merge_collector.fql        OPTIONAL – INSERT INTO: windowed
merge, fresh trace via

ISOTOPE_MERGE_TRACE (fan-in, so it
cannot continue a trace)
  81_merge_edge_markers.fql      OPTIONAL – INSERT INTO: one merge
edge per contributing record,

labelled with the same derived ID
via ISOTOPE_MERGE_TRACE_ID
  99_teardown.fql              DROP TABLE/VIEW/FUNCTION (companion
to cp-flink-reports-down)

```

The three **OPTIONAL** files are applied only when the reports application is started with **--merge-provenance** (**MERGE_PROVENANCE=true scripts/deploy-cmf-flink-reports.sh up**); CCAF's equivalent is **terraform apply -var enable_merge_provenance=true**. Off, none of them is parsed and no extra topic is created. See [docs/flink-collector.md §2.4](#).

CCAF runs the same seven reports plus the collector, but the SQL lives inline as **confluent_flink_statement** resources in [terraform/setup-confluent-flink.tf](#) — 24 statements: ALTER (×4) + raw view + typed produce view + typed consume view + 7 sinks + **STUCK_TRACE_PTF**, **LATENCY_PERCENTILES** and **ISOTOPE_APPEND_HOP** drop+register (×2 each = 6) + the collector sink and its two header ALTERs + **1 statement set holding all 8 INSERTs**. A normal apply creates 28,

because the `ISOTOPE_MERGE_TRACE / ISOTOPE_MERGE_TRACE_ID` drop+register statements (×2 each = 4) live in `terraform/setup-ccaf-merge-provenance.tf` and are *not* gated — they register regardless of `var.enable_merge_provenance`, and only the 5 statements that *call* them are gated. The JAR is uploaded as a `confluent_flink_artifact` and referenced via `USING JAR 'confluent-artifact://<id>'`.

Why one statement rather than eight. CCAF's `EXECUTE STATEMENT SET BEGIN ... END` runs multiple INSERTs as a single optimized statement, and Confluent documents it for exactly this case — INSERTs that read the same table or share intermediate results, which all eight do. It mirrors what CP already does through `IsotopeReportsJob`'s `StatementSet`, and it matters for cost: every separate CCAF statement carries its own 1-CFU compute-pool floor, so eight statements meant eight floors and a saturated pool. The tradeoff is one failure domain — a fault in any INSERT stops them all — which is the same property CP's single job has.

4.0 Wire-format detail (CP only)

Window columns ride as `BIGINT` epoch millis on the wire, not `TIMESTAMP_LTZ`. Flink 2.1.2's `avro-confluent` schema-derivation path raises `UnsupportedOperationException: Unsupported to derive Schema for type: TIMESTAMP_LTZ(3)` for that type. We cast in the INSERT `(UNIX_TIMESTAMP(CAST(window_start AS STRING)) * 1000)` so the on-wire schema is plain Avro `long`. Consumers rehydrate via `T0_TIMESTAMP_LTZ(window_start, 3)`. The CCAF Protobuf sinks don't hit this — `proto-registry` handles `TIMESTAMP_LTZ` directly.

5.0 PTF JAR

The single shadow JAR `ptf/build/libs/isotope-flink-udf.jar` (produced by `./gradlew :ptf:shadowJar`) carries all five JAR-backed functions under the `ai.signalroom.kafka.isotope.flink` package:

- `LatencyPercentilesPTF` (registered as `LATENCY_PERCENTILES`) — T-Digest p50/p95/p99 via per-window state + event-time timers.
- `StuckTracePTF` — per-trace state + event-time timer that emits alerts for traces idle ≥60s.
- `IsotopeAppendHop` (registered as `ISOTOPE_APPEND_HOP`) — a `ScalarFunction`, not a PTF. The two above are *interpreters*, reading headers an interceptor already wrote; this one is a *collector*, appending a Flink hop and rewriting the headers on the way out.
- `IsotopeMergeTrace / IsotopeMergeTraceId` (registered as `ISOTOPE_MERGE_TRACE / ISOTOPE_MERGE_TRACE_ID`) — the optional fan-in collector. `ISOTOPE_APPEND_HOP` continues an inbound trace, which is only truthful 1:1; these mint a **fresh** trace for a merged record and label its merge edges with the same deterministically derived ID. Registered on both runtimes regardless of the feature flag — registration is inert until a statement calls it.

All five register on both runtimes; only the `CREATE FUNCTION ... USING JAR ...` clause differs (`file://` path on CP, `confluent-artifact://` reference on CCAF).

One packaging trap. `Isotope` declares `fromHeaders(org.apache.kafka.common.header.Headers)`. `IsotopeAppendHop` never calls it, but the JVM resolves that descriptor when linking the class, so `Headers` must be present at *runtime*. `compileOnly` is not enough: CCAF's function classloader does not expose `kafka-clients`, and the statement fails on its first row — long after `CREATE FUNCTION` reported success — with `UDF`

invocation error: ... `org.apache.kafka.common.header.Headers`. The JAR therefore bundles `kafka-clients` (`transitive = false`), relocates `org.apache.kafka`, and ships only `common/header`: 2 classes rather than 4,152. `minimize()` cannot substitute — a type named only in a method descriptor is invisible to reachability analysis and gets stripped.

6.0 Operations

6.1 CP Flink on Minikube

```
make cp-flink-up                # cert-manager → CFK Flink Operator → MinIO
→ CMF 2.4 → env → app image
make kafka-pf-up                # localhost:30092 → Kafka, localhost:8081 → SR
make cp-flink-reports-up        # build app JAR → upload as cmf:// artifact →
deploy the CMF Application
make cp-flink-reports-down      # delete the CMF Application + artifact + sink
topics
make cp-flink-down              # tear down the application, CMF, MinIO,
operator, cert-manager
```

The 7 reports run as a single Flink 2.1 CMF **Application** (`isotope-reports`, entry point `IsotopeReportsJob`) — visible in CMF and Control Center's Flink tab.

6.2 CCAF

```
make cc-flink-reports-up  CONFLUENT_API_KEY=... CONFLUENT_API_SECRET=...
                           # terraform apply: env + cluster + topics +
compute pool + artifact + 28 statements
                           # also regenerates terraform/terraform.png via
`terraform graph | dot`
```

or

```
make cc-flink-reports-up  CONFLUENT_API_KEY=... CONFLUENT_API_SECRET=...
ENABLE_TRACE_RCA=true RCA_MODEL_API_KEY=... RCA_MODEL_PROVIDER=...
RCA_MODEL_VERSION=... RCA_MODEL_ENDPOINT=...
                           # terraform apply: env + cluster + topics +
compute pool + artifact + 31 statements (28 + the 3 AI root-cause analysis
statements)
                           # also regenerates terraform/terraform.png via
`terraform graph | dot`
```

```
source scripts/cc-cli-env.sh          # exports BOOTSTRAP / SR_URL /
KAFKA_KEY / KAFKA_SECRET / SR_KEY / SR_SECRET / JAAS
scripts/cc-app-run.sh send orders.placed order-intake-service 'hello' #
drives traffic with the SASL config
```

```
make cc-flink-reports-down CONFLUENT_API_KEY=... CONFLUENT_API_SECRET=...  
                                # terraform destroy: deletes the environment and  
everything in it
```

See the [root README §3.3 "Flink SQL reports on Confluent Cloud for Apache Flink \(CCAF\)"](#) for the full CCAF walkthrough, including the multi-window sustained-traffic pattern required to see tumbling-window aggregates emit.